From Datalog to FLIX: A Declarative Language for Fixed Points on Lattices

Magnus Madsen

University of Waterloo, Canada
mmadsen@uwaterloo.ca

Ming-Ho Yee

University of Waterloo, Canada
ming-ho.yee@uwaterloo.ca

Ondřej Lhoták

University of Waterloo, Canada
olhotak@uwaterloo.ca

Abstract

We present FLIX, a declarative programming language for specifying and solving least fixed point problems, particularly static program analyses. FLIX is inspired by Datalog and extends it with lattices and monotone functions. Using FLIX, implementors of static analyses can express a broader range of analyses than is currently possible in pure Datalog, while retaining its familiar rule-based syntax.

We define a model-theoretic semantics of FLIX as a natural extension of the Datalog semantics. This semantics captures the declarative meaning of FLIX programs without imposing any specific evaluation strategy. An efficient strategy is semi-naïve evaluation which we adapt for FLIX. We have implemented a compiler and runtime for FLIX, and used it to express several well-known static analyses, including the IFDS and IDE algorithms. The declarative nature of FLIX clearly exposes the similarity between these two algorithms.

Categories and Subject Descriptors F3.2 [*Semantics of Programming Languages*]: Program Analysis

General Terms Logic Programming, Static Analysis

Keywords logic programming, static analysis, Datalog

1. Introduction

Least fixed point problems are ubiquitous in mathematics and computer science, significantly in programming languages, and particularly in program analysis. Given a monotone function F on a lattice, the goal is to find the least x for which $F(x) = x$. At the lowest and most general level, a program is a function F that instructs a machine how to change its overall state at each computation step. A static analysis computes an abstract state $\hat{x}$ that over-approximates all possible concrete states that a program can reach. Every sound approximation must satisfy $\hat{F}(\hat{x}) \sqsubseteq \hat{x}$, where $\hat{F}$ is an abstraction of the concrete transformation function F, since if a state in $\hat{x}$ can be reached by a computation, then so can a state in $\hat{F}(\hat{x})$. The least $\hat{x}$ satisfying this property can be computed by starting from the least element $\bot$ and iteratively applying $\hat{F}$ until the fixed point is reached [15, 35].

Static analyzers, which involve fixed-point computations, are complex pieces of software often implemented in general-purpose languages such as C++ or Java. The many mutual dependencies, imposed by the fixed-point problem, are typically expressed using a complex arrangement of worklists. The decision of how to structure the worklists is global, so these large analyzers become difficult to restructure and modify. It also becomes difficult to understand precisely the analysis problem that the implementation is actually solving, and to assure oneself that the implementation is correct. Moreover, the complexity of the dependencies between inter-related sub-analyses often leads analysis implementors to sacrifice precision. For example, some interprocedural analysis frameworks use a call graph precomputed with conservative assumptions to compute dataflow information that would enable a more precise call graph.

To overcome these difficulties, some analysis designers have turned to Datalog [7, 9, 63]. A Datalog program is a set of rules and its solution is the minimal model that satisfies those rules. An important benefit of a Datalog program is its modularity: because the rules are declarative, individual sub-analyses can be easily composed by taking the union of their rules, and the Datalog solver takes care of the mutual dependencies automatically. Thus, it is easy to understand an analysis by understanding its components individually. Correctness of each component implies correctness of the overall analysis. Furthermore, Datalog solvers implement many important optimizations, such as index selection, query planning, and parallel execution [3, 27, 28, 58]. In a hand-crafted static analyzer, each of these optimizations must be implemented manually.

However, Datalog has important limitations that restrict its applicability to program analysis:

Lattices. Datalog is inherently limited to rules on relations, i.e. powersets of tuples, but common static analyses operate over a wide variety of other lattices. Some simple lattices can be embedded in powersets, but at a very high computational cost, and more interesting lattices cannot be encoded at all. For example, we can embed the constant propagation lattice over a finite domain in the following way: $\bot$ is represented by the empty set, each constant is represented by a singleton set, and $\top$ is represented by any set that contains a specially designated $\top$ element. We then add a rule that adds the $\top$ element to every set of two or more elements. However, this $\top$ rule cannot prevent the Datalog program from processing the original non-singleton, non-$\top$ sets. We get the worst of both worlds: the precision is the same as with the constant propagation lattice, but the computational cost is the same as with the much more expensive arbitrary-sets-of-constants lattice. Moreover, when the domain of the constants is infinite, such as the integers, the lattice cannot be encoded at all.

Functions. A related issue is that functions can only be expressed as tabulated relations in Datalog. This can be cumbersome and slow, and functions with an infinite domain or codomain cannot be expressed at all. Consequently, most operations on lattices cannot be expressed. Returning to the constant propagation lattice, even if the lattice itself could somehow be expressed in Datalog, there would be no way to express abstract addition or multiplication on its elements. Moreover, the lack of functions, as well as compound datatypes, means that even a simple context-sensitive analysis such as *k*-CFA cannot be expressed.

Interoperability. A practical limitation of most Datalog solvers is that they do not interact well with existing infrastructure. For instance, static analyzers are often forced to serialize their facts to a file, externally execute the solver, and then read the solution back from a file. Besides additional overhead, this requires a tedious mapping between the analysis and the Datalog solver format.

To overcome these limitations, we propose FLIX, a new declarative language for fixed-point problems. FLIX is a rule-based language inspired by Datalog and extended with lattices and monotone functions.

In summary, this paper makes the following contributions:

- We present FLIX, a new programming language inspired by Datalog and extended with lattices and functions.
- We define a model-theoretic semantics of FLIX as a natural generalization of the Datalog semantics.
- We show how several well-known static analyses, which are inexpressible in Datalog, can be expressed in FLIX.

- We present declarative formulations of the IFDS and IDE algorithms in FLIX, which clearly and concisely demonstrate that IDE is a generalization of IFDS.
- We discuss our implementation of a preliminary FLIX solver and its relation to semi-naïve evaluation. We experimentally compare its performance to hand-crafted static analyzers and the DLV Datalog solver.

2. Motivation

2.1 Points-to Analysis with Datalog

Undoubtedly, the "killer-app" for Datalog has been points-to analysis of object-oriented programs [4, 7, 30, 37, 54, 55, 63]. This involves the specification and computation of sophisticated whole-program subset-based points-to analysis [1, 31]. Consider the following Java fragment:

```
1  ClassA o1 = new ClassA()  // object A
2  ClassB o2 = new ClassB()  // object B
3  ClassB o3 = o2;
4  o2.f = o1;
5  Object r = o3.f;           // Q: What is r?
```

A points-to analysis for this program can help answer questions such as: "to what object can the local variable r point?" In this case, r can point to object A since: (i) variable o1 points to object A, (ii) variable o2 points to object B, (iii) variable o3 points to the value of variable o2 which is object B, (iv) object A is written to field f of object B due to *i* and *ii*, (v) the value of r is object A since variable o3 points to object B due to *iii*, and the value of field f is A due to *iv*.

Unfortunately, such complicated reasoning is necessary due to the multiple recursive dependencies in the dataflow of an object-oriented program: specifically, that local variable information depends on heap information which itself depends on local variable information. Datalog provides an elegant formalism to express such recursive rules. Figure 1 shows a points-to analysis for a minimal object-oriented language. This formulation has just four rules, but is sufficient to capture the entirety of the above reasoning.

The program defines six relations: the four input relations New, Assign, Load, and Store, and the two derived relations VarPointsTo and HeapPointsTo which hold the solution. To analyze the code fragment above, we represent it as a set of *facts* that are fed to a Datalog solver:

```
New("o1", "A").
New("o2", "B").
Assign("o3", "o2").
Store("o2", "f", "o1").
Load("r", "o3", "f").
```

Running the solver infers a solution containing the fact Var-PointsTo("r", "A"), as expected. As an example of how to understand a Datalog program, consider the second rule in Figure 1. This rule states that if there is an assignment from v2 to v1 and v2 points to some object h2, then v1 points to h2. Furthermore, notice the mutual recursion in the rules for VarPointsTo and HeapPointsTo in Figure 1.

```
VarPointsTo(v1, h1) :- New(v1, h1).
VarPointsTo(v1, h2) :- Assign(v1, v2),
                       VarPointsTo(v2, h2).
VarPointsTo(v1, h2) :- Load(v1, v2, f),
                       VarPointsTo(v2, h1),
                       HeapPointsTo(h1, f, h2).
HeapPointsTo(h1, f, h2) :- Store(v1, f, v2),
                       VarPointsTo(v1, h1),
                       VarPointsTo(v2, h2).
```

Figure 1. A field-sensitive subset-based points-to analysis for an object-oriented programming language, e.g. Java.

2.2 Points-to and Dataflow Analysis with FLIX

The analysis presented in the previous section is sufficient if we are only interested in points-to information. However, many static analysis clients require additional information. Consider a client that wants to discover division-by-zero errors. This analysis requires both points-to and dataflow analysis to determine if the denominator in a division is possibly zero. We can use a constant propagation analysis or interval analysis to discover this information. However, as argued in the introduction, such analyses are not expressible in Datalog, but they are expressible in FLIX.

In this section, we show how to develop such an analysis, but for the sake of exposition, we use the parity lattice instead of the constant propagation or interval lattices. As a reminder, the parity lattice tracks whether numbers are odd or even. Figure 2 shows a FLIX program that extends the points-to analysis with a dataflow analysis.

The FLIX language is composed of two parts: a pure functional programming language for specifying lattices, their associated operations, and monotone functions over their elements, combined with a logic language for expressing rules over relations and lattices. Bringing everything together is a set of declarations to specify the names, arities, and types of functions, relations, and lattices. The syntax of FLIX is inspired by Scala and Datalog.

We briefly walk through the program in Figure 2.

The *enum definition* on lines 5–9 defines a tagged union of the elements of the parity lattice. In FLIX, references to a tag must be prefixed by the enum name, thus `Parity.Odd` refers to the odd element of the parity enum.

The *function definitions* on lines 13–20 and 23–24 each define a named function, the type of its arguments, its return type, and its expression body. FLIX functions are used to define the components of a lattice and to express monotone filter and transfer functions. Here, the `leq` function defines the partial order of parity lattice elements. FLIX functions are expressed using a small, pure functional language.

The *lattice definition* on lines 28–29 associates a 5-tuple $(\bot, \top, \sqsubseteq, \sqcup, \sqcap)$ of lattice components with a type, where $\bot$ is the bottom element, $\top$ is the top element, $\sqsubseteq$ is the partial order, $\sqcup$ is the least upper bound, and $\sqcap$ is the greatest lower bound. A lattice definition is a built-in mechanism for defining what is essentially an instance of a type class [62].

```
1   // an almost complete Flix program.
2
3   // an enum definition that defines
4   // the elements of the parity lattice.
5   enum Parity {
6         case Top,
7     case Even, case Odd,
8         case Bot
9   }
10
11  // a function definition that defines
12  // the partial order of the parity lattice.
13  def leq(e1: Parity, e2: Parity): Bool =
14    match (e1, e2) with {
15      case (Parity.Bot, _)         => true
16      case (Parity.Even, Parity.Even) => true
17      case (Parity.Odd, Parity.Odd)   => true
18      case (_, Parity.Top)         => true
19      case _                       => false
20    }
21
22  // additional lattice definitions ...
23  def lub(e1: Parity, e2: Parity): Parity = ...
24  def glb(e1: Parity, e2: Parity): Parity = ...
25
26  // association of the lattice operations
27  // with the parity type.
28  let Parity<> = (Parity.Bot, Parity.Top,
29                  leq, lub, glb);
30
31  // monotone filter and transfer functions ...
32  def isMaybeZero(e: Parity): Bool = ...
33  def sum(e1: Parity, e2: Parity): Parity = ...
34
35  // declaration of relations ...
36  rel Load(var: Str, base: Str, field: Str);
37  rel VarPointsTo(var: Str, obj: Str);
38  // additional declarations ...
39
40  // declaration of lattices ...
41  lat IntVar(var: Str, Parity<>);
42  lat IntField(var: Str, field: Str, Parity<>);
43  // additional declarations ...
44
45  // VarPointsTo and HeapPointsTo rules ...
46
47  // additional dataflow analysis rules ...
48
49  IntVar(v, i) :- Int(v, i).
50  IntVar(v, i) :- Assign(v, v2), IntVar(v2, i).
51  IntVar(v, i) :- Load(v, v2, f),
52                  VarPointsTo(v2, h),
53                  IntField(h, f, i).
54  IntField(h, f, i) :- Store(v1, f, v2),
55                  VarPointsTo(v1, h),
56                  IntVar(v2, i).
57
58  // rule for addition of parity elements.
59  IntVar(r, sum(i1, i2)) :- AddExp(r, v1, v2),
60                  IntVar(v1, i1).
61                  IntVar(v2, i2).
62
63  // rule for potential division-by-zero errors.
64  ArithmeticError(r) :- DivExp(r, v1, v2),
65                  IntVar(v2, i2),
66                  isMaybeZero(i2).
```

Figure 2. A subset-based, field-sensitive points-to analysis combined with a dataflow analysis expressed in FLIX.

Like Haskell, FLIX allows only one type class instance per type. The definition assumes that the supplied functions satisfy the properties of a complete lattice; otherwise, the semantics of the FLIX program is undefined. The notation `Parity<>` is used to distinguish the parity type from the parity type class instance.

The *relation declaration* on line 36 (and 37) defines a relation by specifying its name and its attributes (columns), together with their names and types. The *lattice declaration* on line 41 (and 42) defines a lattice by specifying its name and its attributes (columns), together with their names and types. The last attribute of a lattice declaration must have a type equipped with a lattice. Intuitively, the `IntVar` lattice is the map lattice from strings to elements of the parity lattice.

A *rule definition* in FLIX is similar to a Datalog rule, but is more expressive. First, a FLIX rule may contain function applications in the last term of the head predicate of a rule. For example, consider line 59:

```
IntVar(v, sum(v1, v2)) :- ...
```

Here, the `sum` function is used to compute the abstract sum of two parity lattice elements. The function must be monotone in its arguments. Second, a FLIX rule may use a filter function which is monotone over the booleans. For example, consider lines 64–66:

```
ArithmeticError(r) :- DivExp(r, v1, v2),
                      IntVar(v2, i2),
                      isMaybeZero(i2).
```

Here, the filter function `isMaybeZero` selects only those parity elements i2 that may be zero, i.e. the `Parity.Even` and `Parity.Top` elements.

This example demonstrates three key features of FLIX: the ability to express lattices, monotone filter functions that select a subset of lattice elements, and monotone transfer functions that express mappings between lattice elements. These features go beyond what is possible in Datalog and give FLIX its expressive power.

2.3 Design Choice: Language or Framework?

A natural question to ask is why implement FLIX as a programming language instead of as a framework? First, we want the logic language and the functional language to receive equal treatment. If we had embedded FLIX in a functional language, expressing rules would have suffered, and vice versa. Second, we want full control over the language to ensure that the choices of evaluation strategy and data structures are up to the implementation.

At the same time, we want to ensure interoperability with the JVM, providing access to the existing ecosystem of static analysis frontends and tools. FLIX allows programmers to access JVM types and methods, and use them in their definitions of lattices and functions. Furthermore, FLIX provides an API for accessing computed solutions.

3. Semantics

In this section, we present the model-theoretic semantics of Datalog, and then extend that semantics to FLIX. To do so, we abstract away some details of the full FLIX language.

3.1 Model-Theoretic Semantics of Datalog

Syntax. A Datalog program P is a set of rules of the form $A_0 \Leftarrow A_1, \ldots, A_n$ where A_0 is the *head* of the rule, $A_1, \ldots, A_n$ is the *body* of the rule, and each A_i is an atom. A fact is a rule with an empty body. An atom has the form $p(t_1, \ldots, t_n)$, where p is a predicate symbol and $t_1, \ldots, t_n$ are terms. A term is either a variable x or a constant value v. Values are typically primitive integers and strings. Datalog has no operations on primitive values nor any compound datatypes. Figure 3 shows the syntax of Datalog programs.

Herbrand Universe and Base. We now briefly describe the *Herbrand interpretation* of a Datalog program. This is a simple and elegant approach, commonly used to define the meaning of Datalog programs [9]. The *Herbrand universe $\mathcal{U}$* of a Datalog program P is the set of all possible ground terms. A *ground term* is a non-variable term, i.e. a constant value appearing somewhere in the program P. The Herbrand universe is finite since the program P contains a finite number of constants. The *Herbrand base $\mathcal{B}$* of a Datalog program P is the set of all ground atoms. A *ground atom* is a predicate symbol that occurs in P with its arguments drawn from the Herbrand universe. Note that the Herbrand base respects the arity of the predicates. The Herbrand base is also finite. As an example, given the Datalog program

$$A(1). \qquad B(2,3). \qquad A(x) : - B(x, \_).$$

the Herbrand universe $\mathcal{U}$ is

$$\mathcal{U} = \{1, 2, 3\}$$

and the Herbrand base $\mathcal{B}$ is

$$\mathcal{B} = \begin{cases} A(1), & A(2), & A(3), \\ B(1,1), & B(1,2), & B(1,3), \\ B(2,1), & B(2,2), & B(2,3), \\ B(3,1), & B(3,2), & B(3,3) \end{cases}$$

Intuitively, the Herbrand base can be thought of as the set of "all possible facts."

Interpretations and Models. An *interpretation I* of a Datalog program P is a subset of the Herbrand base $\mathcal{B}$. A ground atom A is true w.r.t. an interpretation I if $A \in I$. A conjunction of atoms $A_1, \ldots, A_n$ is true w.r.t. an interpretation if each atom is true in the interpretation. A ground rule is true if either the body conjunction is false or the head is true. A ground rule is a rule where all atoms are ground.

A *model M* of a Datalog program P is an interpretation, i.e. a subset of the Herbrand base $\mathcal{B}$, that makes each ground

$$
\begin{array}{rcl}
P \in \textit{Prog} & ::= & R_1, \ldots, R_n \\
R \in \textit{Rules} & ::= & A_0 \Leftarrow A_1, \ldots, A_n \\
A \in \textit{Atoms} & = & p(t_1, \ldots, t_n) \\
t \in \textit{Terms} & = & x \mid v \\
p \in \textit{Predicates} & = & \text{is a finite set of predicate symbols.} \\
x \in \textit{Variables} & = & \text{is a finite set of variables.} \\
v \in \textit{Values} & = & \text{is a finite set of values.}
\end{array}
$$

Figure 3. Grammar of Datalog programs.

instance of each rule in P true. A ground instance of a rule is obtained by replacing every variable in a rule with a term from the Herbrand universe. A model M is *minimal* if there is no other model M' such that $M' \subset M$. Consider the following interpretations of the previous program:

$$
\begin{aligned}
I_1 &= \{A(1)\} \\
I_2 &= \{A(1), B(2,3)\} \\
I_3 &= \{A(1), A(2), A(3), B(2,3)\} \\
I_4 &= \{A(1), A(2), B(2,3)\}
\end{aligned}
$$

Here, I_1 and I_2 are not models because they do not make each ground rule instance true. I_3 is a model, but is not minimal as evidenced by the true minimal model I_4.

Computing the Minimal Model. The model-theoretic semantics does not tell us how to compute the minimal model. This is desirable because it separates the definition of *what* the solution is from *how* to compute it. Thus, different Datalog solvers can use different evaluation strategies while still agreeing on the solution.

The model-theoretic semantics does provide some insight into how a solver can be implemented. Let I be an interpretation of a Datalog program P. We define the *immediate consequence operator* T_p of I as the head atoms of each ground rule instance that is satisfied by I. In other words, we can think of I as the current set of facts, and $T_p(I)$ as the set of facts that can be derived from I and P in "one-step." The minimal model of P can then be computed by repeated iteration of $T_p(I)$ starting from the empty set of facts, i.e. it is the least fixed point of $T_p^\infty(\varnothing)$. This strategy is called "naïve" evaluation since it re-evaluates every rule whenever a new fact is inferred. A better strategy is discussed in Section 3.7.

3.2 From Datalog to FLIX

We now extend the Datalog semantics to FLIX.

A *complete lattice* ℓ is a 6-tuple $\ell = (E, \bot, \top, \sqsubseteq, \sqcup, \sqcap)$, where E is a set of elements, $\bot \in E$ is the least element, $\top \in E$ is the greatest element, $\sqsubseteq$ is the partial order on E, $\sqcup$ is the least upper bound, and $\sqcap$ is the greatest lower bound.

We want to allow static analysis implementors to express their own lattices and operations on them. In order to do that, we extend the values of Datalog with enums (tagged unions),

tuples, and sets. Furthermore, we add a pure functional programming language. For the purpose of this paper, this language can be thought of as a minimal lambda calculus, but the full FLIX language is richer. It includes algebraic data types, pattern matching, collections (e.g. lists, sets, and maps), and a static type system.

We make six changes to the Datalog semantics to extend it with lattices. In brief, the steps are:

1. Associate every predicate symbol with a lattice.

2. Extend the Herbrand universe with lattice elements.

3. Partition the Herbrand base into cells and introduce a complete lattice on all ground atoms in each cell.

4. Introduce *compactness* of an interpretation.

5. Extend the definition of a model to incorporate the partial order on ground atoms.

6. Introduce a partial order on models and update the definition of minimality.

We now describe each step in greater detail.

First, we change the grammar of Datalog programs to associate every predicate symbol with a lattice ℓ. Thus, for example, a fact and a rule now look like

$$ A_\ell(\texttt{Odd}). \qquad B_\ell(x) :- A_\ell(x). $$

where ℓ refers to a lattice, in this case, the parity lattice.

Second, we change the definition of the *Herbrand universe* $\mathcal{U}$ to include all possible ground terms $\mathcal{T}$ and lattice elements $\mathcal{E}$ of the FLIX program P. The *Herbrand base* $\mathcal{B}$ remains the same, but uses the new definition of the Herbrand universe $\mathcal{U}$. Notice that the Herbrand universe (and consequently the Herbrand base) may be infinite if a lattice has infinitely many elements.

Third, we partition the Herbrand base such that two ground atoms $A = p_\ell(v_1, \ldots, v_n)$ and $B = p'_{\ell'}(v'_1, \ldots, v'_m)$ are in the same cell S if they have the same predicate symbol (i.e. $p_\ell = p'_{\ell'}$), the same number of terms (i.e. $n = m$), and the first $n - 1$ terms are equal (i.e. $v_1 = v'_1 \wedge \cdots \wedge v_{n-1} = v'_{m-1}$). Notice that all predicate symbols in the same cell have the same associated lattice ℓ. For each cell S, we introduce a complete lattice $L_S = (S, \bot_S, \top_S, \sqsubseteq_S, \sqcup_S, \sqcap_S)$. Given two ground atoms $A = p_\ell(v_1, \ldots, v_n)$ and $B = p_\ell(v'_1, \ldots, v'_n)$, we define their partial order as follows: if $n = 1$ then $A \sqsubseteq_S B$ when $v_1 \sqsubseteq v'_1$. Otherwise, if $n > 1$ then $A \sqsubseteq_S B$ when $v_1 = v'_1 \wedge \cdots \wedge v_{n-1} = v'_{n-1}$ and $v_n \sqsubseteq v'_n$. In words, in order for $A \sqsubseteq_S B$, the first $n - 1$ components of A and B must match exactly, and the last component of A must be less than or equal to the last component of B according to the partial order defined for the lattice ℓ. The definition of the remaining lattice components is straightforward. The intuition is that unary predicates correspond to a single lattice element whereas multi-ary predicates correspond to a single tuple from a map lattice. To ensure a finite number of cells, we restrict the terms that can appear in a ground atom. Recall

that the Herbrand universe $\mathcal{U} = \mathcal{T} \cup \mathcal{E}$ consists of ground terms $\mathcal{T}$ and lattice elements $\mathcal{E}$. We require that in every ground atom $p_\ell(v_1, \ldots, v_n)$, the values $v_1, \ldots, v_{n-1}$ must be drawn from $\mathcal{T}$, and v_n must be drawn from $\mathcal{E}$. Since $\mathcal{T}$ consists of terms that syntactically appear in the program P, it is finite and consequently the number of cells is finite.

Fourth, we introduce the notion of compactness. An interpretation I is *compact* iff every cell S in the partition of I has one unique element.

Fifth, an interpretation is still a subset of the Herbrand base, but we change the definition of when an interpretation is a model. Specifically, a ground atom A is true w.r.t. an interpretation I if $\exists A' \in I$ such that $A \sqsubseteq_S A'$, where A and A' are in the same cell S. As before, a conjunction of atoms $A_1, \ldots, A_n$ is true w.r.t. an interpretation if each atom is true in the interpretation. Finally, a ground rule is true if either the body conjunction is false or the head is true. A model M of P is then an interpretation that makes each ground instance of each rule in P true.

Sixth, we define a partial order $\sqsubseteq_M$ on models. Given two models M_1 and M_2, we say that M_1 is less than or equal to M_2 if for every ground atom $A_1 \in M_1$, belonging to cell S, there is a ground atom $A_2 \in M_2$, also belonging to S, such that $A_1 \sqsubseteq_S A_2$. A model M is *minimal* if it is compact and there is no other model M' such that $M' \sqsubseteq_M M$.

Example. The FLIX program defined over the parity lattice $\ell = (\{\bot, \top, \texttt{Even}, \texttt{Odd}\}, \bot, \top, \sqsubseteq, \sqcup, \sqcap)$ with the facts

$$A_\ell(\texttt{Even}). \quad A_\ell(\texttt{Odd}). \quad B_\ell(\texttt{Odd}).$$

has the Herbrand universe

$$\mathcal{U} = \{\bot, \top, \texttt{Even}, \texttt{Odd}\}$$

and the Herbrand base

$$\mathcal{B} = \begin{Bmatrix} A_\ell(\bot), & A_\ell(\texttt{Even}), & A_\ell(\texttt{Odd}), & A_\ell(\top), \\ B_\ell(\bot), & B_\ell(\texttt{Even}), & B_\ell(\texttt{Odd}), & B_\ell(\top) \end{Bmatrix}$$

An interpretation of the program is a subset of $\mathcal{B}$, e.g.:

$$
\begin{aligned}
I_1 &= \{A_\ell(\top)\} \\
I_2 &= \{A_\ell(\top), B_\ell(\bot)\} \\
I_3 &= \{A_\ell(\top), B_\ell(\texttt{Odd}), B_\ell(\top)\} \\
I_4 &= \{A_\ell(\texttt{Even}), A_\ell(\texttt{Odd}), B_\ell(\texttt{Odd})\} \\
I_5 &= \{A_\ell(\top), B_\ell(\top)\} \\
I_6 &= \{A_\ell(\top), B_\ell(\texttt{Odd})\}
\end{aligned}
$$

The interpretations I_1 and I_2 are not models of the program since neither makes $B_\ell(\texttt{Odd})$ true. I_3 and I_4 are models, but they are not compact. I_5 is a compact model, but it is not minimal as evidenced by the true minimal model I_6.

Example. The FLIX program defined over the sign lattice $\ell = (\{\bot, \top, \texttt{Neg}, \texttt{Zer}, \texttt{Pos}\}, \bot, \top, \sqsubseteq, \sqcup, \sqcap)$ with the facts

$$A_\ell(1, \texttt{Pos}). \quad A_\ell(2, \texttt{Pos}). \quad A_\ell(2, \texttt{Neg}).$$

has the Herbrand universe

$$\mathcal{U} = \{1, 2, \bot, \top, \texttt{Neg}, \texttt{Zer}, \texttt{Pos}\}$$

and the Herbrand base

$$\mathcal{B} = \begin{Bmatrix} A_\ell(1, \bot), & A_\ell(1, \texttt{Neg}), & A_\ell(1, \texttt{Zer}), & \ldots \\ A_\ell(2, \bot), & A_\ell(2, \texttt{Neg}), & A_\ell(2, \texttt{Zer}), & \ldots \end{Bmatrix}$$

An interpretation of the program is a subset of $\mathcal{B}$, e.g.:

$$
\begin{aligned}
I_1 &= \{A_\ell(1, \top)\} \\
I_2 &= \{A_\ell(1, \texttt{Pos}), A_\ell(1, \texttt{Neg}), A_\ell(2, \top)\} \\
I_3 &= \{A_\ell(1, \top), A_\ell(2, \top)\} \\
I_4 &= \{A_\ell(1, \texttt{Pos}), A_\ell(2, \top)\}
\end{aligned}
$$

Here, I_2, I_3 and I_4 are models, I_3 and I_4 are compact, but only I_4 is minimal.

Least Upper and Greatest Lower Bounds. To clarify how lattices are used in FLIX, suppose that we have the two facts:

```
A(Odd).          B(Even).
```

Then the FLIX program with the *two* rules

```
R(x) :- A(x).    R(x) :- B(x).
```

has $R(\top)$ in the minimal model since $\top$ is the only element greater than or equal to both `Odd` and `Even`. On the other hand, the FLIX program with the *one* rule

```
R(x) :- A(x), B(x).
```

has $R(\bot)$ in the minimal model since $\bot$ is the only element less than or equal to both `Odd` and `Even`.

Computing the Minimal Model. We have described how to compute the minimal model of a Datalog program as the least fixed point of the *immediate consequence operator* T_p. We can adopt a similar strategy for FLIX programs with one change: the domain of the operator is restricted to compact interpretations, and since the set of one-step derivable facts may not be compact, we must compute the least upper bound in every cell of the interpretation before we apply T_p again.

Ullman [59, Chapter 3] presents a naïve evaluation algorithm for Datalog with a detailed proof that it terminates and computes the minimal model. The proof relies on two key properties: monotonicity of the immediate consequence operator and finiteness of ascending chains of interpretations. By being careful to maintain monotonicity in defining the FLIX version of the immediate consequence operator, and by insisting that the FLIX lattices be of finite height, we can apply the same proof to a naïve evaluation algorithm for FLIX.

3.3 The Full FLIX Language

The FLIX language is richer than the core formalism presented in the previous subsection. In this subsection, we describe two important extensions: monotone filter functions and monotone transfer functions. Without these features, programming in FLIX would be severely limited. Monotonicity of these functions is necessary to maintain the monotonicity of the immediate consequence operator, which in turn is necessary for the correctness of the naïve evaluation algorithm.

Monotone Filter Functions. A *monotone filter function* is a function from one or more lattice elements to `true` or `false`, and is monotone when the booleans are ordered `false < true`. For example, assume that `isMaybeZero` is a function over the parity lattice elements. Then the rule

```
ArithmeticError(r) :−
    DivExp(r, v1, v2),
    IntVar(v2, i2),
    isMaybeZero(i2).
```

captures a possible division-by-zero error for result variable `r` when there is an expression `r = v1 / v2`, the value of variable `v2` is the parity lattice element `i2`, and `i2` is possibly zero according to the filter function. The important point here is that the `DivExp` relation and `IntVar` lattice are both explicitly represented and tabulated, whereas `isMaybeZero` is a reference to a function expression, i.e. a piece of code. Thus, filter functions may operate over domains that are too difficult or even impossible to tabulate.

We add filter functions to the model-theoretic semantics with the following changes:

1. Extend the definition of a rule to allow filter function applications $f_i(\dots)$ in addition to atoms.

2. Extend the definition of when a ground rule is true. A conjunction of atoms $A_1, \dots, A_n$ and filter function applications $f_1, \dots, f_m$ is true w.r.t. an interpretation if each atom is true in the interpretation, and if each filter function evaluates to true. A ground rule is true if either the body conjunction is false or the head is true.

Monotone Functions. A *monotone function* is a function from one or more lattice elements to a lattice element, which is order-preserving. These functions are used to implement "transfer" functions in FLIX. For example, in the rule

```
IntVar(r, sum(i1, i2)) :−
    AddExp(r, v1, v2),
    IntVar(v1, i1),
    IntVar(v2, i2).
```

the monotone function `sum` computes the abstract addition of two parity lattice elements `i1` and `i2`. We highlight two important design choices here.

First, we require such transfer functions to be strict and monotone. Strictness ensures that when a function is applied to $\bot$ it returns $\bot$. Monotonicity requires that the function is order-preserving and is necessary to ensure that the immediate consequence operator is also monotone.

Second, we only allow non-filter functions to appear in the last term of the head predicate of a rule. This ensures that the implementation can evaluate the rule body first, in any order, and then evaluate the head. If non-filter functions could appear in the body, then one could write

```
R(z) :− A(x), B(y, f(x)), C(z, g(y)).
R(x) :− A(y, f(x)), B(z, g(y)), C(x, h(z)).
```

where the first rule implicitly enforces an evaluation order[1]. Moreover, it is not clear how to evaluate the second rule. Thus, we disallow non-filter functions in rule bodies.

We add monotone "transfer" functions to the model-theoretic semantics with the following changes:

1. Extend the definition of a rule to allow the last term of a head predicate to contain function applications $f_i(\dots)$.

2. Extend the definition of a ground rule instance to allow function applications in the last term of a head predicate.

3. Extend the definition of when a ground atom is true. A ground atom A with function applications is true iff after all function applications are evaluated, the resulting ground atom is true.

3.4 Compositionality of FLIX Programs

An important property of Datalog programs is compositionality: given two disjoint Datalog programs P_1 and P_2, i.e. two programs that share no predicate symbols, the union of their rules is a Datalog program P. The model of P is the union of the models of P_1 and P_2. FLIX retains this property, allowing composition of analyses. For example, given a constant propagation analysis and a reachability analysis, we can combine the two to obtain a conditional constant propagation analysis. We share information between the two analyses by introducing the shared predicates `isReachable`(s), `isTrue`(s), and `isFalse`(s). The constant propagation analysis infers facts for the `isTrue`(s) and `isFalse`(s) predicates and uses the `isReachable`(s) predicate, while the reachability analysis does the opposite.

In abstract interpretation terminology, the compositionality discussed above is known as the *direct product* [14]. Given two analyses, e.g. sign and parity, the direct product corresponds to running each analysis independently. No information is shared and neither lattice is used to refine the other. For example, the element $(\texttt{Zer}, \texttt{Odd})$ in the Cartesian product `Sign` $\times$ `Parity` does not correspond to any concrete value and could be replaced by $(\bot, \bot)$. The *reduced product* uses this idea to refine and share information between lattices [11, 14]. The *logical product* is more precise than the reduced product, and under certain conditions can be constructed automatically [29]. FLIX provides the direct product automatically, but the reduced and logical products must be implemented manually.

---

[1] Since the functions `f` and `g` are opaque and may not have an inverse.

3.5 Negation, Stable Models, and Stratification

Pure Datalog does not allow negation: every atom in the body of a rule must appear unnegated. Many proposals have been made to extend logic and Datalog programs with negation [20, 21, 23, 26, 36]. A fundamental challenge is how to define the semantics of programs such as:

```
A(x) :- !B(x).        B(x) :- !A(x).
```

To overcome these issues, several solutions including stable models [25], well-founded semantics [61], and many-valued logic [23] have been developed. A common solution is to restrict the use of negation to so-called *stratifiable* programs [2]. A program is stratifiable if it satisfies a simple syntactic property which ensures that no negative cycles occur in the program. FLIX currently does not support any form of negation, but it is something we plan to add.

3.6 The Theoretical Expressive Power of Datalog

We have previously stated that certain static analyses are "inexpressible" in pure Datalog. Strictly speaking, this is not true: a well-known result from logic and complexity theory shows that the data complexity of Datalog is PTIME-complete [17]. The *data complexity* of a program is the complexity when the input is restricted to facts (i.e. the rules are kept fixed) and PTIME is the class of all polynomial time algorithms. Intuitively, any algorithm which runs in polynomial time can be *encoded* as a Datalog program. Naturally, this includes polynomial time static analyses.

However, this encoding simulates a Random Access Machine that tabulates all steps taken by the algorithm, and furthermore, requires operations such as integer arithmetic to be represented as finite relations. While this is sufficient for an existence proof, such tabulation is completely impractical for an implementation.

3.7 The Semi-naïve Evaluation Strategy

The model-theoretic semantics of FLIX describes the structure of the minimal model of a program, but not how to compute it. This is a good thing, since it gives the fixed-point solver maximum freedom in the choice of data structures and evaluation strategy. As explained before, the model-theoretic semantics inspires a naïve evaluation strategy based on the immediate consequence operator. The idea is to repeatedly re-evaluate every rule while maintaining a monotonically growing set of facts. While this strategy is simple and correct, it is hopelessly inefficient.

A better strategy, known as *semi-naïve evaluation*, tracks the dependencies between predicates and predicate symbols appearing in rule bodies. Under this evaluation strategy, whenever a new fact for a predicate p is inferred, only the rules containing p in their body are re-evaluated, and only for the new fact. For example, in the FLIX program

```
SelfLoop(x) :- Edge(x, x).
Path(x, y)  :- Edge(x, y).
Path(x, z)  :- Path(x, y), Edge(y, z).
```

if the fact `Path(1, 2)` is inferred, then *only* the third rule is re-evaluated under the environment $\{x \mapsto 1, \ y \mapsto 2\}$. The first and second rules are *not* re-evaluated. With naïve evaluation, all three rules would be re-evaluated.

We use a variant of semi-naïve evaluation, adapted for FLIX by taking the compactness requirement into account. For example, in the FLIX program

```
A(Odd).
B(Even).
A(x) :- B(x).
R(x) :- isMaybeZero(x), A(x).
```

we initially infer the two facts $A(\text{Odd})$ and $B(\text{Even})$, which cause evaluation of the two rules. This infers the new fact $A(\text{Even})$. Next, we must re-evaluate the third rule, but we must *not* use the environment $\{x \mapsto \text{Even}\}$ since this breaks the compactness requirement. Instead, we must compute the least upper bound of $A(\text{Odd}) \sqcup A(\text{Even}) = A(\top)$, and re-evaluate the third rule under the environment $\{x \mapsto \top\}$. This gives the correct minimal model M where $R(\top) \in M$.

In addition to naïve evaluation, Ullman [59, Chapter 3] also defines and proves correctness of a semi-naïve evaluation algorithm for Datalog. The semi-naïve algorithm maintains a so-called incremental relation for each predicate. Whenever the algorithm evaluates a rule with head predicate p_i to yield a new relation P_i', the incremental relation is computed as the set difference $\Delta P_i = P_i' \setminus P_i$, where P_i is the old relation for the predicate p_i. The incremental relation contains only the facts that were newly computed in P_i', that were not already contained in the old relation P_i.

Then, the process of evaluating a rule is modified to the following procedure. Letting n be the number of atoms in the body of the rule, the rule is evaluated n times. Each time, one of the n atoms is instantiated using the incremental relation ΔP_i corresponding to the predicate specified by the atom. All other atoms of the body of the rule are instantiated using their corresponding old relations P_i. Informally, this ensures that each newly-inferred fact in ΔP_i is considered with all of the existing facts in the old relations for the other atoms. Ullman proves that each such incremental evaluation step yields the same resulting relation for the head predicate as the corresponding full evaluation step in the naïve algorithm. Furthermore, he proves inductively that every iteration of the semi-naïve algorithm infers the same facts as the corresponding iteration of the naïve algorithm (but using less work), and therefore the final outputs of the two algorithms are the same.

To adapt semi-naïve evaluation to FLIX, we must first adapt the definition of the incremental relation. In Datalog, $P_i' \supseteq P_i$ by the monotonicity of rule evaluation, so $P_i' = \Delta P_i \cup P_i$. This latter property is important because each step of semi-naïve evaluation can be informally thought of as determining $\Delta P_i \cup P_i$, while the corresponding step of naïve evaluation computes P_i'. However, in FLIX, it is not generally true that $P_i' \supseteq P_i$ because rule evaluation applies a least upper bound in each cell S to obtain a compact

relation P_i'. Instead of $P_i' \sqsupseteq P_i$, FLIX ensures that $P_i' \sqsupseteq P_i$. An alternative definition of the incremental relation ΔP_i is therefore necessary so that $P_i' = \Delta P_i \sqcup P_i$, analogously to the Datalog property $P_i' = \Delta P_i \cup P_i$. We have defined a set of facts P to be compact if it contains one ground atom for each cell S; let us denote this ground atom $ga(P, S)$. When P_i' and P_i are the relations for a predicate before and after evaluating a rule, we define their incremental relation $\Delta P_i = \{ ga(P_i', S) \mid S \in \text{cells} \wedge ga(P_i', S) \sqsupsetneq ga(P_i, S) \}$. In words, the incremental relation ΔP_i contains every ground atom from P_i' which is strictly greater than the ground atom for the same cell in P_i. Since $P_i' \sqsupseteq P_i$, this definition of ΔP_i does imply the desired property $P_i' = \Delta P_i \sqcup P_i$.

The process of evaluating a rule can then be incrementalized in the same way as for Datalog, in that the rule is evaluated as many times as there are atoms in its body, and each time, one of the atoms is instantiated with the FLIX version of the incremental relation ΔP_i. The resulting relation for the head predicate is then made compact by computing a least upper bound for each cell S. We can show that this incremental rule evaluation step yields the same resulting relation for the head predicate as the corresponding full evaluation step, in the same way as the analogous incremental evaluation step does in Datalog. Therefore, the inductive proof of the equivalence of outputs of naïve and semi-naïve evaluation for Datalog then applies analogously to FLIX.

4. Evaluation

We have implemented a compiler and runtime for FLIX. The entire implementation is roughly 30,000 lines of Scala code. The toolchain includes a parser, a type checker, an interpreter, an indexed database, and a semi-naïve fixed-point solver. The source code is freely available on GitHub.[2]

We evaluate the usefulness of FLIX by implementing three existing static analyses: the Strong Update analysis [39] and the IFDS and IDE algorithms [52, 53].

4.1 The Strong Update Analysis

The Strong Update analysis is a points-to analysis for C programs that propagates singleton points-to sets flow sensitively and larger sets flow insensitively. The paper presents the analysis as a set of constraints (Figure 7 in that paper) and then as an imperative algorithm (Figure 9 in that paper). The latter is implemented in C++ and evaluated [39]. The FLIX implementation shown in Figure 4 follows the constraint specification of the analysis directly: there is a one-to-one correspondence between the FLIX rules and the constraints from Figure 7 of the Strong Update paper. The PtSU function is defined as:

$$ptsu[\bar{\ell}](a) \triangleq \begin{cases} su[\bar{\ell}](a) & \text{if } su[\bar{\ell}](a) \neq \top \\ pt(a) & \text{if } su[\bar{\ell}](a) = \top \end{cases}$$

The monotonicity of this definition depends on the unstated fact that $su[\bar{\ell}](a) \subseteq pt(a)$ when $su[\bar{\ell}](a)$ is a singleton set.

```
enum SULattice {
  case Top,
  case Single(Str),
  case Bottom
}

def filter(t: SULattice, b: Str): Bool =
  match t with {
    case SULattice.Bottom     => false
    case SULattice.Single(p)  => b == p
    case SULattice.Top        => true
  }

Pt(p, a) :- AddrOf(p, a).
Pt(p, a) :- Copy(p, q), Pt(q, a).
Pt(p, b) :- Load(l, p, q), Pt(q, a), PtSU(l, a, b).
PtH(a, b) :- Store(l, p, q), Pt(p, a), Pt(q, b).

SUBefore(l2, a, t) :-
    CFG(l1, l2),
    SUAfter(l1, a, t).
SUAfter(l, a, t) :-
    SUBefore(l, a, t),
    Preserve(l, a).
SUAfter(l, a, SULattice.Single(b)) :-
    Store(l, p, q),
    Pt(p, a), Pt(q, b).
PtSU(l, a, b) :-
    PtH(a, b),
    SUBefore(l, a, t),
    filter(t, b).
```

Figure 4. FLIX version of the Strong Update analysis [39]. SUBefore and SUAfter is the information before and after label l, written $ptsu[\bar{l}](a)$ and $ptsu[\underline{l}](a)$ in the paper. Preserve is the complement of the Kill set.

In the corresponding FLIX rule, monotonicity is explicit: the rule first selects all $b \in pt(a)$, and then uses the function filter to reject all elements other than p in the case that $su[\bar{\ell}](a)$ is the singleton $\{p\}$.

4.2 IFDS

IFDS is a framework that can be instantiated to solve a large class of interprocedural context-sensitive dataflow analyses, the interprocedural finite distributive subset analyses [52]. The paper that defines the framework presents it as a one-page algorithm in pseudocode that contains many worklist updates and implicit quantifications. Anecdotally, many people find the algorithm difficult to understand, and checking its correctness requires a long proof.

Figure 5 declaratively specifies the desired properties of an IFDS solution. It is also a set of FLIX rules that can be executed to compute the solution. The rules compute a set of *path edges*, each leading from a data point d_1 at the start of a procedure to a data point d_3 at an instruction m within the procedure, and a set of *summary edges* that summarize the transfer function of each call to a procedure. To implement a specific dataflow analysis, one must provide the transfer functions of that analysis. These are specified to FLIX in the form of three functions eshIntra (intraprocedural),

```
PathEdge(d1, m, d3) :-
    CFG(n, m),
    PathEdge(d1, n, d2),
    d3 <- eshIntra(n, d2).
PathEdge(d1, m, d3) :-
    CFG(n, m),
    PathEdge(d1, n, d2),
    SummaryEdge(n, d2, d3).
PathEdge(d3, start, d3) :-
    PathEdge(d1, call, d2),
    CallGraph(call, target),
    EshCallStart(call, d2, target, d3),
    StartNode(target, start).
SummaryEdge(call, d4, d5) :-
    CallGraph(call, target),
    StartNode(target, start),
    EndNode(target, end),
    EshCallStart(call, d4, target, d1),
    PathEdge(d1, end, d2),
    d5 <- eshEndReturn(target, d2, call).

EshCallStart(call, d, target, d2) :-
    PathEdge(_, call, d),
    CallGraph(call, target),
    d2 <- eshCallStart(call, d, target).

Result(n, d2) :-
    PathEdge(_, n, d2).
```

Figure 5. FLIX implementation of the IFDS analysis [52].

```
JumpFn(d1, m, d3, comp(long, short)) :-
    CFG(n, m),
    JumpFn(d1, n, d2, long),
    (d3, short) <- eshIntra(n, d2).
JumpFn(d1, m, d3, comp(caller, summary)) :-
    CFG(n, m),
    JumpFn(d1, n, d2, caller),
    SummaryFn(n, d2, d3, summary).
JumpFn(d3, start, d3, identity()) :-
    JumpFn(d1, call, d2, _),
    CallGraph(call, target),
    EshCallStart(call, d2, target, d3, _),
    StartNode(target, start),
SummaryFn(call, d4, d5, comp(comp(cs, se), er)) :-
    CallGraph(call, target),
    StartNode(target, start),
    EndNode(target, end),
    EshCallStart(call, d4, target, d1, cs),
    JumpFn(d1, end, d2, se),
    (d5, er) <- eshEndReturn(target, d2, call).

EshCallStart(call, d, target, d2, cs) :-
    JumpFn(_, call, d, _),
    CallGraph(call, target),
    (d2, cs) <- eshCallStart(call, d, target).

InProc(p, start) :- StartNode(p, start).
InProc(p, m) :- InProc(p, n), CFG(n, m).

Result(n, d, apply(fn, vp)) :-
    ResultProc(proc, dp, vp),
    InProc(proc, n),
    JumpFn(dp, n, d, fn).

ResultProc(proc, dp, apply(cs, v)) :-
    Result(call, d, v),
    EshCallStart(call, d, proc, dp, cs).
```

Figure 6. FLIX implementation of the IDE analysis [53].

eshCallStart (call site to start node), and eshEndReturn (end node to return site). The special arrow syntax in the first rule d3 <- eshIntra(n, d2) binds d3 to each element in the set returned by the call.

It is essential that the transfer functions be specified as functions; they cannot be tabulated and given as input to any solver in the form of relations. The reason is that the overall goal of the IFDS algorithm is to compute a set of reachable pairs (n, d). If we knew the set of arguments (n, d) for which eshIntra needs to be tabulated (the set of all arguments with which the analysis will call eshIntra), then we would already know which pairs are reachable and there would be no need to run the IFDS analysis. Alternatively, tabulating eshIntra for all possible pairs (n, d) would be much more costly than performing the IFDS algorithm itself, which calls eshIntra only on the much smaller subset of pairs (n, d) that are reachable.

Another important detail not discussed in the IFDS paper is that the algorithm applies not only the eshCallStart function (in the third rule) but also its inverse (in the fourth, SummaryEdge rule). Since computing the inverse is usually impractical, any implementation of the algorithm must tabulate the function for the arguments on which it is called in the forward direction. The FLIX program explicitly tabulates the function in the EshCallStart relation. As a result, the relation can be consulted in both directions in the third and fourth rule, and this can be written declaratively as it is in the original formulation of the IFDS algorithm.

Both of these issues are discussed by Naeem et al. [48].

4.3 IDE

IDE is a more general framework that can be instantiated to solve a larger class of interprocedural context-sensitive dataflow analyses, the interprocedural distributive environment analyses [53]. The original presentation of IDE as an imperative algorithm requires two pages. Conceptually, the IDE framework is a direct extension of the IFDS framework, but that is not obvious at all from the worklist-based algorithmic specifications in the two papers. The IDE framework computes the same edges as IFDS, but each edge is decorated with a representation of a so-called micro-function. This correspondence between IFDS and IDE is immediately clear from the declarative specification of IDE shown in Figure 6. Notice that the rules mirror those of the IFDS implementation (with the names PathEdge and SummaryEdge replaced with JumpFn and SummaryFn to match the terminology used in the IFDS and IDE papers). In the IDE algorithm, each rule has just one additional component corresponding to the micro-function on each edge. The first, second, and fourth rules use a FLIX function comp in the head to compute the composition of micro-functions.

```
def comp(t1: Transformer, t2: Transformer): Transformer = match (t1, t2) with {
  case (_, BotTransformer)                          => BotTransformer
  case (BotTransformer, NonBotTransformer(a, b, Value.Bot))    => BotTransformer
  case (BotTransformer, NonBotTransformer(a, b, Value.Cst(k))) => NonBotTransformer(0, k, Value.Cst(k))
  case (BotTransformer, NonBotTransformer(a, b, Value.Top))    => NonBotTransformer(0, 0, Value.Top)
  case (NonBotTransformer(a2, b2, c2), NonBotTransformer(a1, b1, c1)) =>
    NonBotTransformer(a1 * a2, (a1 * b2) + b1, lub(sum(prod(c2, a1), b1), c1))
}
```

Figure 7. Micro-function join operation for the example IDE analysis [53].

To instantiate the IDE framework with a specific analysis, one must not only implement the transfer functions `eshIntra`, `eshCallStart`, and `eshEndReturn`, but also specify two lattices: the value lattice V that is the domain and range of each micro-function, and the micro-function lattice F that efficiently represents certain functions from $V \rightarrow V$. Thus, the formulation of the IDE algorithm requires both functions and lattices.

The IDE paper uses a running example of a linear constant propagation analysis in which V is the constant propagation lattice. The elements of the micro-function lattice F are $\lambda l.\bot$ and functions of the form $\lambda l.(a \times l + b) \sqcup c$, where a and b are integers and c is an element of the constant propagation lattice. Figure 7 shows the FLIX implementation of the micro-function composition operation `comp` that is called from the rules. The functions that implement $\sqsubseteq$ and $\sqcup$ on the lattice of micro-functions have a similar structure.

4.4 Shortest Paths

We have shown how to express several static analyses in FLIX, but FLIX is applicable to other types of fixed-point problems. For example, to compute all-pairs shortest paths, let $(\mathcal{N}, \infty, 0, \geqslant, \texttt{min}, \texttt{max})$ be a lattice over the natural numbers. Then we can compute the shortest paths as follows:

```
Dist(y, d + c) :- Dist(x, d), Edge(x, y, c).
```

4.5 Performance of the Current FLIX Solver

We have not yet looked closely at the performance of the FLIX solver. Instead, we have focused on language design and semantics. For that reason, the solver contains many inefficiencies that can be overcome with additional engineering. For example, primitive values (e.g. integers) are represented as (boxed) Java objects, functions (including the partial order and least upper bound) are evaluated using an AST-based interpreter, rules are always evaluated left-to-right instead of using a cost-plan, relations are represented as hash maps but some would be more efficiently represented as dense arrays, and our index selection strategy is not optimal.

Table 1 compares the performance of three implementations of the Strong Update analysis on the benchmark programs evaluated in the Strong Update paper. First, we implemented the Strong Update analysis in Datalog by embedding the Strong Update lattice within the relational powerset lattice as described in the introduction. We used the well-known DLV Datalog solver [38] to run this implementation.

The Datalog version of the analysis did not scale beyond the 458.sjeng benchmark (13.9 kSLOC), which it analyzed in 425 seconds. The FLIX formulation of the analysis analyzed the same benchmark in 27 seconds, and was able to scale up to the 300.twolf benchmark (20.5 kSLOC). We confirmed that both implementations compute the same results. The C++ implementation in LLVM from the original Strong Update paper is still much faster. This is partly, but not entirely, due to the constant overheads of the current implementation of FLIX. In addition, the C++ implementation uses a clever data structure to implement a map from abstract objects. Thanks to subtle properties of the Strong Update analysis, the data structure can avoid explicitly representing the objects whose corresponding lattice value is either $\top$ or $\bot$ in the common case. FLIX, on the other hand, explicitly represents all objects whose corresponding lattice value is either a singleton or $\top$. The number of objects whose lattice value is $\top$ is very large, which accounts for much of the performance difference.

We also evaluated the performance of the IFDS analysis described previously. As a concrete example of a specific IFDS analysis, we selected the object abstraction from the multi-object typestate analysis of Naeem et al. [47]. In the FLIX implementation, we provided an interface that enables monotone transfer functions to be implemented using Java or Scala code, rather than the FLIX functional language. This made it possible for the declarative FLIX IFDS program to call the same implementations of the transfer functions that were evaluated in the original typestate analysis paper.

Specifically, our evaluation compares two implementations of the object abstraction analysis. The baseline is the complete implementation that was used in the original paper. It includes a hand-coded imperative Scala implementation of the IFDS algorithm, and Scala implementations of the IFDS functions that instantiate the IFDS framework to compute the object abstraction analysis. The FLIX implementation includes the declarative IFDS formulation, which is instantiated with the same Scala implementations of the object abstraction analysis functions. Thus, the evaluation compares the hand-coded imperative implementation of IFDS with the declarative FLIX implementation. We verified that both implementations produce the same outputs.

The running times of the two implementations on the six DaCapo benchmarks [6] that were evaluated in the original paper are shown in Table 2. In general, the current FLIX im-

	Benchmark		**DLV**		**Flix**		**C++**
Program	kSLOC	Input Facts	Memory (MB)	Time (s)	Memory (MB)	Time (s)	Time (s)
470.lbm	1.2	1,205	17	1.8	142	0.9	0.05
181.mcf	2.5	3,377	114	30.9	650	3.0	0.08
429.mcf	2.7	3,392	115	31.8	630	3.1	0.09
256.bzip2	4.7	5,017	49	5.9	244	1.8	0.09
462.libquantum	4.4	6,196	215	32.3	877	5.2	0.14
164.gzip	8.6	9,259	463	133.4	1,271	9.4	0.14
401.bzip2	8.3	11,844	1,100	696.4	2,264	17.5	0.30
458.sjeng	13.9	20,154	1,077	424.8	3,107	27.1	0.27
433.milc	15.0	22,147	-	timeout	3,846	88.6	0.45
175.vpr	17.8	25,977	-	-	4,039	99.7	0.54
186.crafty	21.2	32,189	-	-	3,556	73.0	0.41
197.parser	11.4	32,606	-	-	5,104	663.5	0.92
482.sphinx3	25.1	42,736	-	-	6,767	399.7	1.06
300.twolf	20.5	44,041	-	-	5,273	222.8	1.25
456.hmmer	36.0	68,384	-	-	-	timeout	2.22
464.h264ref	51.6	89,898	-	-	-	-	3.41
	seven more benchmarks						

Table 1. Summary of performance results for the Strong Update analysis. Timeout means more than 15 minutes.

Program	**Scala** Time (s)	**Flix** Time (s)	**Slowdown**
luindex	133.6	366.7	2.7x
antlr	176.7	437.3	2.5x
hsqldb	187.4	469.2	2.5x
bloat	203.5	584.1	2.9x
pmd	247.7	680.1	2.7x
jython	4,614.7	14,344.8	3.1x

Table 2. Summary of performance results for IFDS.

plementation of the IFDS algorithm is about 3x slower than the imperative Scala implementation. Importantly, the performance of FLIX scales with the imperative implementation.

5. Related Work

Static Analysis Frameworks. Many static analysis frameworks have been proposed over the years. The Program Analysis Generator (PAG) generates C code for static analyzers from specifications of lattices and descriptions of transfer functions [44]. The TJ Watson WALA library is a static analysis library written in Java [22]. WALA includes implementations of many common static analyses such as points-to analysis, class hierarchy analysis, and the IFDS algorithm [52]. Soot is a Java bytecode analysis framework that has seen wide use in compilation and as a frontend for other static analyses [60]. Hoopl is a generic dataflow analysis and transformation framework written in Haskell [51]. Frama-C is a source code analysis platform for C programs written in OCaml [16]. A weakness of all these frameworks is the implicit assumptions that they make about the analysis, for instance, the existence of a control-flow graph or other intermediate representation, and the choice of flow-sensitivity,

context-sensitivity, and memory abstraction. FLIX, like Datalog, makes no such assumptions and is applicable to least fixed point problems in general.

Datalog. Datalog has roots in the database community as a general-purpose query language, but has found uses in many areas of computer science. A comprehensive introduction to Datalog is given by Ceri et al. [9]. A distilled version is given by Huang et al. who argue that interest in Datalog is re-emerging [32]. Further evidence of this is provided by the Datalog 2.0 Workshop [19].

The use of negation in logic programs, such as Datalog and Prolog, has long been studied and there is extensive literature on the subject [2, 20, 21, 23, 25, 26, 36]. The theoretical complexity of logic programs, including Datalog, has also been studied extensively [17].

Datalog has been used in several static analysis tools. codeQuest is a tool for querying various aspects of the source code of a program [30]. Binary decision diagrams (BDDs) have been used to implement efficient points-to analyses specified as relations, both using a custom relational language and Datalog [40, 63]. The Doop framework is a precise and scalable context-sensitive points-to analysis for Java specified in Datalog [7, 54, 55].

FLIX has some similarities to Bloom, a programming language designed to ensure consistency of distributed programs [13]. Bloom, like FLIX, takes inspiration from Datalog and adds support for lattices and monotone functions, but for different reasons. In Bloom, the purpose is to ensure *confluence*, i.e. that regardless of the order in which messages are received over the network, the same result is computed. Unlike Datalog and FLIX, where the user is interested in some minimal model, Bloom programs are intended to run continuously as network services.

Logic Programming. Prolog [10, 56] is a Turing-complete logic programming language related to Datalog. Every Datalog program is also a Prolog program; however, Prolog allows constructors, negation, and the cut operator. Datalog programs can be efficiently solved by Prolog engines that use tabulation, e.g. XSB [57]. Such Prolog engines have been used to implement different types of program analyses [18]. FLIX is less expressive than Prolog, but ensures that every program terminates and has a unique minimal model. In future work we plan to compare FLIX to Prolog engines.

Constraint Logic Programming (CLP) schemes extend logic programming with a decidable background theory, such as lists, trees, and linear arithmetic [12, 33, 34, 41]. Intuitively, a CLP program is a set of Horn clauses, each equipped with a formula over the background theory. During evaluation, term unification is augmented with the decision procedure of the theory. CLP has been used in program verification and abstract interpretation [5, 24].

Alternation-free Least Fixed Point (ALFP) logic is an extension of Horn clauses with nested universal and existential quantification, stratified negation, and disjunction [49, 50]. It is more powerful than Datalog, but is still guaranteed to have a minimal solution, provided that the program is stratifiable.

Answer Set Programming (ASP) is a logic tailored to solving NP-hard problems that may involve non-monotonic reasoning [8, 42, 43]. A key development was the introduction of *stable models* (later known as *answer sets*) that allow models to be defined even in the presence of negative cycles in the implication graph [25]. FLIX does not yet support negation, so the challenges of monotonicity are restricted to ensuring that user-defined functions are monotone.

Prolog, CLP, ALFP, ASP, and other logic languages all offer various trade-offs in terms of performance, expressive power and safety (i.e. the existence of least models). FLIX is yet another point in this large design space and is more general than Datalog, but less general than Prolog.

Static Analyzers based on Horn Clauses. Logic programs, such as Prolog and CLP programs, have also been the subject of static analysis. In this line of research, static analyzers such as PLAI [46] compute the least fixed point of a set of Horn clauses over one or more lattices. These lattices and their associated operations are specified as "plug-ins" to the analysis. A key difference between these techniques and FLIX is the choice of evaluation strategy: FLIX uses bottom-up semi-naïve evaluation rather than top-down evaluation with tabulation, as in XSB. These developments have been used to translate object-oriented programs into constraint Horn clauses on which the static analysis is then performed [45].

6.   Conclusion

We have presented FLIX, a declarative programming language for expressing and solving least fixed point problems, particularly static program analyses. FLIX is inspired by Datalog and extends it with lattices and monotone functions. We have defined a model-theoretic semantics for FLIX that is the foundation for any FLIX fixed-point solver, independent of its specific evaluation strategy. We have demonstrated the expressiveness of FLIX by implementing several well-known static analyses, including the IFDS and IDE algorithms. The declarative formulation of these analyses clearly reveals their close relationship. Experimental results show that the current FLIX interpreter is slower than hand-crafted analyzers, but we plan to address this in future work.

7.   Future Work

We briefly outline three directions for future work:

Negation. FLIX does not support negation. We believe it is straightforward to extend the semantics and implementation to support stratified FLIX programs. However, we want to explore whether stratification is the right choice for specifying static analyses. Furthermore, we want to look at potentially interesting connections between negation and lattices.

Safety. Every Datalog program is guaranteed to terminate with the unique minimal model. As discussed previously, a FLIX program is also guaranteed to terminate with the minimal model, provided that every lattice is actually a complete lattice, of finite height, and every function is strict and monotone. Unfortunately, a FLIX programmer may inadvertently violate one or more of the required properties when specifying a lattice or function. We plan to investigate the use of automatic program verification techniques to guarantee that FLIX programs are meaningful.

Performance. As discussed in the evaluation, the current implementation of FLIX has many opportunities for improved performance, such as eliminating boxing, replacing the interpreter with compiled JVM bytecode, and using query planning and better index selection.

Acknowledgment

References

[1] L. O. Andersen. *Program Analysis and Specialization for the C Programming Language*. PhD thesis, University of Copenhagen, 1994.

[2] K. R. Apt, H. A. Blair, and A. Walker. Towards a Theory of Declarative Knowledge. In *Foundations of Deductive Databases and Logic Programming*. 1988. doi: `10.1016/B978-0-934613-40-8.50006-3`.

[3] F. Bancilhon, D. Maier, Y. Sagiv, and J. D. Ullman. Magic Sets and Other Strange Ways to Implement Logic Programs. In *Proc. Principles of Database Systems (PODS)*, 1985. doi: `10.1145/6012.15399`.

[4] W. C. Benton and C. N. Fischer. Interactive, Scalable, Declarative Program Analysis: From Prototype to Implementation.

In *Proc. Principles and Practice of Declarative Programming (PPDP)*, 2007. doi: 10.1145/1273920.1273923.

[5] N. Bjørner, K. McMillan, and A. Rybalchenko. On Solving Universally Quantified Horn Clauses. In *Proc. Symposium on Static Analysis (SAS)*, 2013. doi: 10.1007/978-3-642-38856-9_8.

[6] S. M. Blackburn, R. Garner, C. Hoffman, A. M. Khan, K. S. McKinley, R. Bentzur, A. Diwan, D. Feinberg, D. Frampton, S. Z. Guyer, M. Hirzel, A. Hosking, M. Jump, H. Lee, J. E. B. Moss, A. Phansalkar, D. Stefanović, T. VanDrunen, D. von Dincklage, and B. Wiedermann. The DaCapo benchmarks: Java Benchmarking Development and Analysis. In *Proc. Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*, 2006. doi: 10.1145/1167473.1167488.

[7] M. Bravenboer and Y. Smaragdakis. Strictly Declarative Specification of Sophisticated Points-To Analyses. In *Proc. Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*, 2009. doi: 10.1145/1640089.1640108.

[8] G. Brewka, T. Eiter, and M. Truszczyński. Answer Set Programming at a Glance. *Communications of the ACM*, 2011. doi: 10.1145/2043174.2043195.

[9] S. Ceri, G. Gottlob, and L. Tanca. What You Always Wanted to Know About Datalog (and Never Dared to Ask). *IEEE Transactions on Knowledge and Data Engineering (TKDE)*, 1989. doi: 10.1109/69.43410.

[10] W. Clocksin and C. S. Mellish. *Programming in Prolog*. Springer Berlin Heidelberg, 2003. doi: 10.1007/978-3-642-55481-0.

[11] M. Codish, A. Mulkers, M. Bruynooghe, M. G. De La Banda, and M. Hermenegildo. Improving Abstract Interpretations by Combining Domains. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, 1995. doi: 10.1145/200994.200998.

[12] J. Cohen. Constraint Logic Programming Languages. *Communications of the ACM*, 1990. doi: 10.1145/79204.79209.

[13] N. Conway, W. R. Marczak, P. Alvaro, J. M. Hellerstein, and D. Maier. Logic and Lattices for Distributed Programming. In *Proc. Symposium on Cloud Computing (SoCC)*, 2012. doi: 10.1145/2391229.2391230.

[14] A. Cortesi, G. Costantini, and P. Ferrara. A Survey on Product Operators in Abstract Interpretation. In *Proc. Festschrift for David Schmidt*, 2013. doi: 10.4204/EPTCS.129.19.

[15] P. Cousot and R. Cousot. Abstract Interpretation: A Unified Lattice Model for Static Analysis of Programs by Construction or Approximation of Fixpoints. In *Proc. Principles of Programming Languages (POPL)*, 1977. doi: 10.1145/512950.512973.

[16] P. Cuoq, F. Kirchner, N. Kosmatov, V. Prevosto, J. Signoles, and B. Yakobowski. Frama-C – A Software Analysis Perspective. In *Software Engineering and Formal Methods (SEFM)*, 2012. doi: 10.1007/978-3-642-33826-7_16.

[17] E. Dantsin, T. Eiter, G. Gottlob, and A. Voronkov. Complexity and Expressive Power of Logic Programming. *ACM Computing Surveys (CSUR)*, 2001. doi: 10.1145/502807.502810.

[18] S. Dawson, C. Ramakrishnan, and D. Warren. Practical Program Analysis Using General Purpose Logic Programming Systems – A Case Study. In *Proc. Programming Language Design and Implementation (PLDI)*, 1996. doi: 10.1145/231379.231399.

[19] O. de Moor, G. Gottlob, T. Furche, and A. Sellers, editors. *Datalog Reloaded – First International Workshop, Datalog 2010*, 2011. doi: 10.1007/978-3-642-24206-9.

[20] T. Eiter, G. Gottlob, and H. Mannila. Adding Disjunction to Datalog. In *Proc. Principles of Database Systems (PODS)*, 1994. doi: 10.1145/182591.182639.

[21] T. Eiter, G. Gottlob, and H. Mannila. Disjunctive Datalog. *ACM Transactions on Database Systems (TODS)*, 1997. doi: 10.1145/261124.261126.

[22] S. Fink and J. Dolby. WALA – The TJ Watson Libraries for Analysis, 2012.

[23] M. Fitting. Fixpoint Semantics for Logic Programming a Survey. *Theoretical Computer Science (TCS)*, 2002. doi: 10.1016/S0304-3975(00)00330-3.

[24] C. Flanagan. Automatic Software Model Checking Using CLP. In *Proc. European Symposium on Programming (ESOP)*, 2003. doi: 10.1007/3-540-36575-3_14.

[25] M. Gelfond and V. Lifschitz. The Stable Model Semantics for Logic Programming. In *Proc. International Conference on Logic Programming (ICLP/SLP)*, 1988.

[26] M. Gelfond and V. Lifschitz. Classical Negation in Logic Programs and Disjunctive Databases. *New Generation Computing*, 1991. doi: 10.1007/BF03037169.

[27] G. Graefe. Query Evaluation Techniques for Large Databases. *ACM Computing Surveys (CSUR)*, 1993. doi: 10.1145/152610.152611.

[28] S. Gregory. *Parallel Logic Programming in PARLOG: The Language and its Implementation*. Addison-Wesley, 1987.

[29] S. Gulwani and A. Tiwari. Combining Abstract Interpreters. In *Proc. Programming Language Design and Implementation (PLDI)*, 2006. doi: 10.1145/1133981.1134026.

[30] E. Hajiyev, M. Verbaere, and O. D. Moor. codeQuest: Scalable Source Code Queries with Datalog. In *Proc. European Conference on Object-Oriented Programming (ECOOP)*, 2006. doi: 10.1007/11785477_2.

[31] M. Hind. Pointer Analysis: Haven't We Solved This Problem Yet? In *Proc. Program Analysis for Software Tools and Engineering (PASTE)*, 2001. doi: 10.1145/379605.379665.

[32] S. S. Huang, T. J. Green, and B. T. Loo. Datalog and Emerging Applications: An Interactive Tutorial. In *Proc. Management of Data (SIGMOD)*, 2011. doi: 10.1145/1989323.1989456.

[33] J. Jaffar and J.-L. Lassez. Constraint Logic Programming. In *Proc. Principles of Programming Languages (POPL)*, 1987. doi: 10.1145/41625.41635.

[34] J. Jaffar and M. J. Maher. Constraint Logic Programming: A Survey. *Journal of Logic Programming*, 1994. doi: 10.1016/0743-1066(94)90033-7.

[35] J. B. Kam and J. D. Ullman. Monotone Data Flow Analysis Frameworks. *Acta Informatica*, 1977. doi: 10.1007/BF00290339.

[36] K. Kunen. Negation in Logic Programming. *Journal of Logic Programming*, 1987. doi: `10.1016/0743-1066(87)90007-0`.

[37] M. S. Lam, J. Whaley, V. B. Livshits, M. C. Martin, D. Avots, M. Carbin, and C. Unkel. Context-sensitive Program Analysis as Database Queries. In *Proc. Principles of Database Systems (PODS)*, 2005. doi: `10.1145/1065167.1065169`.

[38] N. Leone, G. Pfeifer, W. Faber, T. Eiter, G. Gottlob, S. Perri, and F. Scarcello. The DLV System for Knowledge Representation and Reasoning. *ACM Transactions on Computational Logic (TOCL)*, 2006. doi: `10.1145/1149114.1149117`.

[39] O. Lhoták and K.-C. A. Chung. Points-To Analysis with Efficient Strong Updates. In *Proc. Principles of Programming Languages (POPL)*, 2011. doi: `10.1145/1925844.1926389`.

[40] O. Lhoták and L. Hendren. Scaling Java Points-To Analysis using Spark. In *Proc. Compiler Construction (CC)*, 2003. doi: `10.1007/3-540-36579-6_12`.

[41] N. Li and J. C. Mitchell. Datalog with Constraints: A Foundation for Trust Management Languages. In *Proc. Practical Aspects of Declarative Languages (PADL)*, 2003. doi: `10.1007/3-540-36388-2_6`.

[42] V. Lifschitz. Answer Set Planning. In *Proc. Logic Programming and Nonmonotonic Reasoning (LPNMR)*, 1999. doi: `10.1007/3-540-46767-X_28`.

[43] V. Lifschitz. What Is Answer Set Programming? In *Proc. Artificial Intelligence (AAAI)*, 2008.

[44] F. Martin. PAG – An Efficient Program Analyzer Generator. *Journal on Software Tools for Technology Transfer (STTT)*, 1998. doi: `10.1007/s100090050017`.

[45] M. Méndez-Lojo, J. Navas, and M. Hermenegildo. A Flexible, (C)LP-Based Approach to the Analysis of Object-Oriented Programs. In *Proc. Logic-Based Program Synthesis and Transformation (LOPSTR)*, 2007. doi: `10.1007/978-3-540-78769-3_11`.

[46] K. Muthukumar and M. Hermenegildo. Compile-time Derivation of Variable Dependency Using Abstract Interpretation. *Journal of Logic Programming*, 1992. doi: `10.1016/0743-1066(92)90035-2`.

[47] N. A. Naeem and O. Lhoták. Typestate-like Analysis of Multiple Interacting Objects. In *Proc. Object-Oriented Programming, Systems, Languages, and Applications (OOPSLA)*, 2008. doi: `10.1145/1449764.1449792`.

[48] N. A. Naeem, O. Lhoták, and J. Rodriguez. Practical Extensions to the IFDS Algorithm. In *Proc. Compiler Construction (CC)*, 2010.

[49] F. Nielson, H. R. Nielson, and H. Seidl. A Succinct Solver for ALFP. *Nordic Journal of Computing (NJC)*, 2002.

[50] F. Nielson, H. R. Nielson, H. Sun, M. Buchholtz, R. R. Hansen, H. Pilegaard, and H. Seidl. The Succinct Solver Suite. In *Proc. Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, 2004. doi: `10.1007/978-3-540-24730-2_21`.

[51] N. Ramsey, J. Dias, and S. P. Jones. Hoopl: A Modular, Reusable Library for Dataflow Analysis and Transformation. In *Proc. Haskell Symposium*, 2010. doi: `10.1145/1863523.1863539`.

[52] T. Reps, S. Horwitz, and M. Sagiv. Precise Interprocedural Dataflow Analysis via Graph Reachability. In *Proc. Principles of Programming Languages (POPL)*, 1995. doi: `10.1145/199448.199462`.

[53] M. Sagiv, T. Reps, and S. Horwitz. Precise Interprocedural Dataflow Analysis with Applications to Constant Propagation. *Theoretical Computer Science (TCS)*, 1996. doi: `10.1016/0304-3975(96)00072-2`.

[54] Y. Smaragdakis and M. Bravenboer. Using Datalog for Fast and Easy Program Analysis. In *Datalog Reloaded*, 2011. doi: `10.1145/1926385.1926390`.

[55] Y. Smaragdakis, M. Bravenboer, and O. Lhoták. Pick Your Contexts Well: Understanding Object-Sensitivity. In *Proc. Principles of Programming Languages (POPL)*, 2011. doi: `10.1145/1925844.1926390`.

[56] L. Sterling and E. Y. Shapiro. *The Art of Prolog: Advanced Programming Techniques*. MIT Press, 1994.

[57] T. Swift and D. S. Warren. XSB: Extending Prolog with Tabled Logic Programming. *Theory and Practice of Logic Programming (TPLP)*, 2012. doi: `10.1017/S1471068411000500`.

[58] J. D. Ullman. *Principles of Database Systems*. Galgotia publications, 1984.

[59] J. D. Ullman. *Principles of Database and Knowledge-Base Systems, Volume I*. Computer Science Press, 1988.

[60] R. Vallée-Rai, P. Co, E. Gagnon, L. Hendren, P. Lam, and V. Sundaresan. Soot – A Java Bytecode Optimization Framework. In *Proc. Centre for Advanced Studies on Collaborative Research (CASCON)*, 1999.

[61] A. Van Gelder, K. A. Ross, and J. S. Schlipf. The Well-Founded Semantics for General Logic Programs. *Journal of the ACM (JACM)*, 1991. doi: `10.1145/116825.116838`.

[62] P. Wadler and S. Blott. How to Make Ad-hoc Polymorphism Less Ad Hoc. In *Proc. Principles of Programming Languages (POPL)*, 1989. doi: `10.1145/75277.75283`.

[63] J. Whaley and M. S. Lam. Cloning-based Context-sensitive Pointer Alias Analysis using Binary Decision Diagrams. In *Proc. Programming Language Design and Implementation (PLDI)*, 2004. doi: `10.1145/996893.996859`.